

R Recap

Mark Dunning

Last modified: 28 Sep 2016

A short introduction to R

Outline

In this session, we review some of the fundamentals of the R language. It should be a useful refresher prior to the intermediate-level Data Analysis and Visualisation using R course.

Topics covered include:-

- Creating variables
- Using Functions
- Vectors
- Data frame
- Subsetting data, the base R way
- Plotting
- Statistical testing
- How to get help

For a more detailed introduction, we suggest the following *free* resources

- Solving Biological Problems with R
- Introduction to Data Science with R
- Coursera course in R
- Beginners Introduction to R Statistical Software
- R programming wiki
- Quick R

R basics

Advantages of R

The R programming language is now recognised beyond the academic community as an effect solution for data analysis and visualisation. Notable users of R include:-

- Facebook,
- google,
- Microsoft (who recently invested in a commerical provider of R)
- The New York Times.
- Buzzfeed use R for some of their serious articles (i.e. not the ones featuring cat pictures), and have made the code publically available
- The New Zealand Tourist Board have R running in the background of their website

The New York Times

Business Computing

WORLD U.S. N.Y. / REGION BUSINESS TECHNOLOGY SCIENCE HEALTH SPORTS OPINION

Search Technology

Go

Inside Technology

Internet Start-Ups Business Computing Companies

Bits
Blo

Data Analysts Captivated by R's Power



Stuart Isett for The New York Times

R first appeared in 1996, when the statistics professors Robert Gentleman, left, and Ross Ihaka released the code as a free software package.

Figure 1:



Key features

- Open-source
- Cross-platform
- Access to existing visualisation / statistical tools
- Flexibility
- Visualisation and interactivity
- Add-ons for many fields of research
- Facilitating ***Reproducible Research***

Two Biostatisticians (later termed ‘*Forensic Bioinformaticians*’) from M.D. Anderson used R extensively during their re-analysis and investigation of a Clinical Prognostication paper from Duke. The subsequent scandal put Reproducible Research at the forefront of everyone’s mind.

Keith Baggerly’s talk on the subject is highly-recommended.

Support for R

- Online forums such as Stack Overflow regularly feature R
- Blogs
- Local user groups
- Documentation via `? or help.start()`
- Documentation for packages is found via the Packages tab in the bottom-right of RStudio.
- Packages analyse all kinds of Genomic data (>800)
- Compulsory documentation (*vignettes*) for each package
- 6-month release cycle
- Course Materials
- Example data and workflows
- Common, re-usable framework and functionality
- Available Support
 - Often you will be able to interact with the package maintainers / developers and other power-users of the project software
- Annual conferences in U.S and Europe
 - The last European conference was in Cambridge

Submit Your

RESEARCH

75 COMMENTS

How Bright Promise in Cancer Testing Fell Apart

By GINA KOLATA JULY 7, 2011

Email

Share

Tweet

Pin

Save

More

When Juliet Jacobs found out she had lung [cancer](#), she was terrified, but realized that her hope lay in getting the best treatment medicine could offer. So she got a second opinion, then a third. In February of 2010, she ended up at [Duke University](#), where she entered a research study whose promise seemed stunning.

Doctors would assess her [tumor](#) cells, looking for gene patterns that would determine which [drugs](#) would best



Keith Baggerly, left, and Kevin Coombes, statisticians at M. D. Anderson Cancer Center, found flaws in research on tumors.

Michael Stravato for The New York Times

Figure 2: duke-scandal



Figure 3:

RStudio

- Rstudio is a free environment for R
- Convenient menus to access scripts, display plots
- Still need to use *command-line* to get things done
- Developed by some of the leading R programmers
- Used by beginners, and experienced users alike

To get started, you will need to install the latest version of R and RStudio Desktop; both of which are *free*.

Once installed, you should be able to launch RStudio by clicking on its icon:-



Figure 4:

The bottom-left with the blinking cursor is the R console > and ready for you to start entering R commands

Getting started



Figure 5:

At a basic level, we can use R as a calculator to compute simple sums with the +, -, * (for multiplication) and / (for division) symbols.

```
2 + 2
```

```
## [1] 4
```

```
2 - 2
```

```
## [1] 0
```

```
4 * 3
```

```
## [1] 12
```

```
10 / 2
```

```
## [1] 5
```

The answer is displayed at the console with a [1] in front of it. The 1 inside the square brackets is a place-holder to signify how many values were in the answer (in this case only one). We will talk about dealing with lists of numbers shortly...

In the case of expressions involving multiple operations, R respects the BODMAS system to decide the order in which operations should be performed.

```
2 + 2 * 3
```

```
## [1] 8
```

```
2 + (2 * 3)
```

```
## [1] 8
```

```
(2 + 2) * 3
```

```
## [1] 12
```

R is capable of more complicated arithmetic such as trigonometry and logarithms; like you would find on a fancy scientific calculator. Of course, R also has a plethora of statistical operations as we will see.

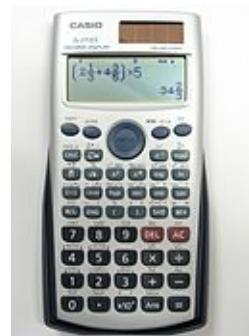


Figure 6:

```
pi
```

```
## [1] 3.141593
```

```
sin(pi/2)
```

```
## [1] 1
```

```
cos(pi)
```

```
## [1] -1
```

```
tan(2)  
  
## [1] -2.18504  
  
log(1)  
  
## [1] 0
```

We can only go so far with performing simple calculations like this. Eventually we will need to store our results for later use. For this, we need to make use of *variables*.

Variables

A variable is a letter or word which takes (or contains) a value. We use the assignment ‘operator’, `<-` to create a variable and store some value in it.

```
x <- 10  
x  
  
## [1] 10  
  
myNumber <- 25  
myNumber  
  
## [1] 25
```

We also can perform arithmetic on variables using functions:

```
sqrt(myNumber)  
  
## [1] 5
```

We can add variables together:

```
x + myNumber  
  
## [1] 35
```

We can change the value of an existing variable:

```
x <- 21  
x  
  
## [1] 21
```

- We can set one variable to equal the value of another variable:

```
x <- myNumber  
x
```

```
## [1] 25
```

- We can modify the contents of a variable:

```
myNumber <- myNumber + sqrt(16)  
myNumber
```

```
## [1] 29
```

When we are feeling lazy we might give our variables short names (`x`, `y`, `i...` etc), but a better practice would be to give them meaningful names. There are some restrictions on creating variable names. They cannot start with a number or contain characters such as `.`, `_`, `'`. Naming variables the same as in-built functions in R, such as `c`, `T`, `mean` should also be avoided.

Naming variables is a matter of taste. Some conventions exist such as a separating words with `-` or using `camelCaps`. Whatever convention you decided, stick with it!

Functions

Functions in R perform operations on **arguments** (the inputs(s) to the function). We have already used:

```
sin(x)
```

this returns the sine of `x`. In this case the function has one argument: `x`. Arguments are always contained in parentheses – curved brackets, `()` – separated by commas.

Arguments can be named or unnamed, but if they are unnamed they must be ordered (we will see later how to find the right order). The names of the arguments are determined by the author of the function and can be found in the help page for the function. When testing code, it is easier and safer to name the arguments. `seq` is a function for generating a numeric sequence *from* and *to* particular numbers. Type `?seq` to get the help page for this function.

```
seq(from = 2, to = 20, by = 4)
```

```
## [1] 2 6 10 14 18
```

```
seq(2, 20, 4)
```

```
## [1] 2 6 10 14 18
```

Arguments can have *default* values, meaning we do not need to specify values for these in order to run the function.

`rnorm` is a function that will generate a series of values from a *normal distribution*. In order to use the function, we need to tell R how many values we want

```
rnorm(n=10)
```

```
## [1] 1.27539694 1.71887371 0.39423972 0.23886899 -0.94258542
## [6] -0.10256703 -0.69723258 1.01811983 -1.21167694 -0.07007651
```

The normal distribution is defined by a *mean* (average) and *standard deviation* (spread). However, in the above example we didn't tell R what mean and standard deviation we wanted. So how does R know what to do? All arguments to a function and their default values are listed in the help page

(*N.B sometimes help pages can describe more than one function*)

```
?rnorm
```

In this case, we see that the defaults for mean and standard deviation are 0 and 1. We can change the function to generate values from a distribution with a different mean and standard deviation using the *mean* and *sd* *arguments*. It is important that we get the spelling of these arguments exactly right, otherwise R will an error message, or (worse?) do something unexpected.

```
rnorm(n=10, mean=2, sd=3)
```

```
## [1] 0.9601574 1.5756687 4.1956100 1.6459811 6.7025385 2.3626814
## [7] -1.0251317 2.3349803 1.3718703 5.8146413
```

```
rnorm(10, 2, 3)
```

```
## [1] 5.5625346 -7.2953643 -1.7354863 0.5143433 5.2175571 3.7066023
## [7] -0.3289767 7.7021639 5.1199219 2.3097185
```

In the examples above, *seq* and *rnorm* were both outputting a series of numbers, which is called a *vector* in R and is the most-fundamental data-type.

Exercise

- How can we create a sequence from 2 to 20 comprised of 5 equally-spaced numbers?
 - check the help page for *seq* `?seq`
- What is the value of *pi* to 3 decimal places?
 - see the help for *round* `?round`

Vectors

- The basic data structure in R is a **vector** – an ordered collection of values.
- R treats even single values as 1-element vectors.
- The function `c` *combines* its arguments into a vector:
- Remember that as `c` is a function, we specify it's arguments in curved brackets(...)

```
x <- c(3,4,5,6)
x
```

```
## [1] 3 4 5 6
```

The `seq` function we saw before was another example of how to create a sequence of values. A useful shortcut is to use the `:` symbol.

```
x <- 3:6
x
```

```
## [1] 3 4 5 6
```

Exercise

- `rep` can be used to replicate values. Construct the following sequences
 - 1 1 1 1 1 2 2 2 2 2
 - 1 2 1 2 1 2 1 2 1 2
 - 1 2 1 2 1 2 1 2 1 2 1
 - (this last sequence has 11 values in it)

The square brackets [] indicate the position within the vector (the *index*). We can extract individual elements by using the [] notation:

```
x[1]
```

```
## [1] 3
```

```
x[4]
```

```
## [1] 6
```

We can even put a vector inside the square brackets: (*vector indexing*)

Exercise

Without using R!

- If $y \leftarrow 2:4$, what would $x[y]$ give?
 - [1] 3 5
 - [1] 2 4
 - [1] 4 5 6

When applying all standard arithmetic operations to vectors, application is element-wise. Thus, we say that R supports *vectorised* operations.

```
x <- 1:10  
y <- x*2  
y
```

```
## [1] 2 4 6 8 10 12 14 16 18 20
```

```
z <- x^2  
x + y
```

```
## [1] 3 6 9 12 15 18 21 24 27 30
```

Vectorised operations are extremely powerful. Operations that would require a *for* loop (or similar) in other languages such as **C**, **Python**, can be performed in a single line of R code.

A vector can also contain text; called a character vector. Such a vector can also be constructed using the **c** function.

```
x <- c("A", "B", "C", "D")
```

The quote marks are crucial. Why?

```
x <- c(A, B, C, D)
```

```
## Error in try(x <- c(A, B, C, D), silent = TRUE) : object 'A' not found
```

Another useful type of data that we will see is the *logical* or *boolean* which can take either the values of TRUE or FALSE

```
x <- c(TRUE, TRUE, FALSE)
```

Logical values are useful when we want to create subsets of our data. We can use *comparison* operators; ==, >, <, != to check if values are equal, greater than, less than, or not equal.

```
x <- c("A", "A", "B", "B", "C")
x == "A"
```

```
## [1] TRUE TRUE FALSE FALSE FALSE
```

```
x != "A"
```

```
## [1] FALSE FALSE TRUE TRUE TRUE
```

```
x <- rnorm(10)
x > 0
```

```
## [1] TRUE TRUE TRUE TRUE FALSE FALSE FALSE FALSE FALSE
```

However, all items in the vector **must** be the same type. If you attempt anything else, R will convert all values to the same (most basic) type.

```
x <- c(1, 2, "three")
x
```

```
## [1] "1"      "2"      "three"
```

Packages in R

So far we have used functions that are available with the *base* distribution of R; the functions you get with a clean install of R. The open-source nature of R encourages others to write their own functions for their particular data-type or analyses.

Packages are distributed through *repositories*. The most-common ones are CRAN and Bioconductor. CRAN alone has many thousands of packages.

The **Packages** tab in the bottom-right panel of RStudio lists all packages that you currently have installed. Clicking on a package name will show a list of functions that available once that package has been loaded. The **library** function is used to load a package and make it's functions / data available in your current R session. *You need to do this every time you load a new RStudio session.*

```
library(beadarray)
```

There are functions for installing packages within R. If your package is part of the main **CRAN** repository, you can use `install.packages`

We will be using the `gapminder` R package in this practical. To install it, we would do.

```
install.packages("gapminder")
```

Bioconductor packages have their own install script, which you can download from the Bioconductor website

```
source("http://www.bioconductor.org/biocLite.R")
biocLite("affy")
```

A package may have several *dependancies*; other R packages from which it uses functions or data types (re-using code from other packages is strongly-encouraged). If this is the case, the other R packages will be located and installed too.

So long as you stick with the same version of R, you won't need to repeat this install process.

Dealing with data

We are going to explore some of the basic features of R using data from the gapminder project, which have been bundled into an R package. These data give various indicator variables for different countries around the world (life expectancy, population and Gross Domestic Product). We have saved these data as a `.csv` file to demonstrate how to import data into R.

You can download these data here. Right-click the link and save to somewhere on your computer that you wish to work from.

The working directory

Like other software (Word, Excel, Photoshop...), R has a default location where it will save files to and import data from. This is known as the *working directory* in R. You can query what R currently considers its working directory by doing:-

```
getwd()
```

N.B. Here, a set of open and closed brackets () is used to call the `getwd` function with no arguments.

We can also list the files in this directory with:-

```
list.files()
```

Any `.csv` file in the working directory can be imported into R by supplying the name of the file to the `read.csv` function and creating a new variable to store the result. A useful sanity check is the `file.exists` function which will print `TRUE` if the file can be found in the working directory.

```
file.exists("gapminder.csv")
```

```
## [1] TRUE
```

If the file we want to read is not in the current working directory, we will have to write the path to the file; either *relevant* to the current working directory (e.g. the directory “up” from the current working directory, or in a sub-folder), or the full path. In an interactive session, you can do use `file.choose` to open a dialogue box. The path to the the file will then be displayed in R.

```
file.choose()
```

Assuming the file can be found, we can use `read.csv` to import. Other functions can be used to read tab-delimited files (`read.delim`) or a generic `read.table` function. A data frame object is created.

```
gapminder <- read.csv("gapminder.csv")
```

The data frame object in R allows us to work with “tabular” data, like we might be used to dealing with in Excel, where our data can be thought of having rows and columns. The values in each column have to all be of the same type (i.e. all numbers or all text).

Exercise

- What are the dimensions of the data frame?
 - What columns are available?
 - HINT: see the `dim`, `ncol`, `nrow` and `colnames` functions
-
-
-

In Rstudio , you can view the contents of the data frame we have just created. This is useful for interactive exploration of the data, but not so useful for automation and scripting and analyses.

```
View(gapminder)
```

We should always check the data frame that we have created. Sometimes R will happily read data using an inappropriate function and create an object without raising an error. However, the data might be unsusable. Consider:-

```
test <- read.delim("gapminder.csv")
head(test)

##           country continent year lifeExp pop gdpPercap
## 1  Afghanistan   Asia 1952  28.801 8425333 779.4453145
## 2  Afghanistan   Asia 1957  30.332 9240934 820.8530296
## 3  Afghanistan   Asia 1962  31.997 10267083 853.10071
## 4  Afghanistan   Asia 1967  34.02  11537966 836.1971382
## 5  Afghanistan   Asia 1972  36.088 13079460 739.9811058
## 6  Afghanistan   Asia 1977  38.438 14880372 786.11336
```

```
dim(test)  
## [1] 1704     1
```

We can access the columns of a data frame by knowing the column name. **TIP** Use auto-complete with the **TAB** key to get the name of the column correct

```
gapminder$country
```

A vector (1-dimensional) is returned, the length of which is the same as the number of rows in the data frame. The vector could be stored as a variable and itself be subset or used in further calculations

The **summary** function is a useful way of summarising the data containing in each column. It will give information about the *type* of data (remember, data frames can have a mixture of numeric and character columns) and also an appropriate summary. For numeric columns, it will report some stats about the distribution of the data. For categorical data, it will report the different *levels*.

```
summary(gapminder)
```

```
##          country      continent       year      lifeExp  
## Afghanistan: 12    Africa :624   Min.   :1952   Min.   :23.60  
## Albania     : 12   Americas:300   1st Qu.:1966   1st Qu.:48.20  
## Algeria     : 12   Asia    :396   Median  :1980   Median  :60.71  
## Angola      : 12   Europe   :360   Mean    :1980   Mean    :59.47  
## Argentina   : 12   Oceania : 24   3rd Qu.:1993   3rd Qu.:70.85  
## Australia   : 12                           Max.   :2007   Max.   :82.60  
## (Other)     :1632  
##          pop           gdpPercap  
## Min.   :6.001e+04   Min.   : 241.2  
## 1st Qu.:2.794e+06   1st Qu.: 1202.1  
## Median :7.024e+06   Median : 3531.8  
## Mean   :2.960e+07   Mean   : 7215.3  
## 3rd Qu.:1.959e+07   3rd Qu.: 9325.5  
## Max.   :1.319e+09   Max.   :113523.1  
##
```

Exercise

- Save the life expectancy and population as variables
 - what is the maximum life expectancy?
 - what is the smallest population?
 - round the life expectancy and populations to the nearest whole numbers
 - HINT:- `min`, `max`, `round`.....
-
-
-

Subsetting

A data frame can be subset using square brackets [] placed after the name of the data frame. As a data frame is a two-dimensional object, you need a *row* and *column* index, or vector indices.

```
gapminder[1,2]
gapminder[2,1]
gapminder[c(1,2,3),1]
gapminder[c(1,2,3),c(1,2,3)]
```

Note that the data frame is not altered we are just seeing what a subset of the data looks like and not changing the underlying data. If we wanted to do this, we would need to create a new variable.

```
gapminder
```

Should we wish to see all rows, or all columns, we can neglect either the row or column index

```
gapminder[1,]
gapminder[,1]
```

Just like subsetting a vector, the indices can be vectors containing multiple values

```
gapminder[1:3,1:2]
gapminder[seq(1,1704,length.out = 10),1:4]
```

A common shortcut is `head` which prints the first six rows of a data frame.

```
head(gapminder)
```

```
##          country continent year lifeExp      pop gdpPercap
## 1 Afghanistan      Asia 1952  28.801  8425333  779.4453
## 2 Afghanistan      Asia 1957  30.332  9240934  820.8530
## 3 Afghanistan      Asia 1962  31.997 10267083  853.1007
## 4 Afghanistan      Asia 1967  34.020 11537966  836.1971
## 5 Afghanistan      Asia 1972  36.088 13079460  739.9811
## 6 Afghanistan      Asia 1977  38.438 14880372  786.1134
```

When subsetting entire rows ***you need to remember the , after the row indices***. If you fail to do so, R may still return a result. However, it probably won't be what you expected. Look what happens if you wanted to the first three rows but typed the following command

```
gapminder[1:3]
```

Rather than selecting rows based on their *numeric* index (as in the previous example) we can use what we call a *logical test*. This is a test that gives either a TRUE or FALSE result. When applied to subsetting, only rows with a TRUE result get returned.

For example we could compare the `lifeExp` variable to 40. The result is a *vector* of TRUE or FALSE; one for each row in the data frame

```
gapminder$lifeExp < 40
```

This R code can be put inside the square brackets to select rows of interest (those observations where the life expectancy variable is less than 40).

```
gapminder[gapminder$lifeExp < 40, ]
```

```
##      country continent year lifeExp      pop gdpPercap
## 1 Afghanistan     Asia 1952  28.801 8425333 779.4453
## 2 Afghanistan     Asia 1957  30.332 9240934 820.8530
## 3 Afghanistan     Asia 1962  31.997 10267083 853.1007
## 4 Afghanistan     Asia 1967  34.020 11537966 836.1971
## 5 Afghanistan     Asia 1972  36.088 13079460 739.9811
## 6 Afghanistan     Asia 1977  38.438 14880372 786.1134
```

The , is important as this tells R to display all columns. If we wanted a subset of the columns we would put their indices after the ,

```
gapminder[gapminder$lifeExp < 40, 1:4]
```

```
##      country continent year lifeExp
## 1 Afghanistan     Asia 1952  28.801
## 2 Afghanistan     Asia 1957  30.332
## 3 Afghanistan     Asia 1962  31.997
## 4 Afghanistan     Asia 1967  34.020
## 5 Afghanistan     Asia 1972  36.088
## 6 Afghanistan     Asia 1977  38.438
```

Testing for equality can be done using ==. This will only give TRUE for entries that are *exactly* the same as the test string.

```
gapminder[gapminder$country == "Zambia", ]
```

```
##      country continent year lifeExp      pop gdpPercap
## 1681 Zambia     Africa 1952  42.038 2672000 1147.389
## 1682 Zambia     Africa 1957  44.077 3016000 1311.957
## 1683 Zambia     Africa 1962  46.023 3421000 1452.726
## 1684 Zambia     Africa 1967  47.768 3900000 1777.077
## 1685 Zambia     Africa 1972  50.107 4506497 1773.498
## 1686 Zambia     Africa 1977  51.386 5216550 1588.688
## 1687 Zambia     Africa 1982  51.821 6100407 1408.679
## 1688 Zambia     Africa 1987  50.821 7272406 1213.315
## 1689 Zambia     Africa 1992  46.100 8381163 1210.885
## 1690 Zambia     Africa 1997  40.238 9417789 1071.354
## 1691 Zambia     Africa 2002  39.193 10595811 1071.614
## 1692 Zambia     Africa 2007  42.384 11746035 1271.212
```

N.B. For partial matches, the grep function and / or *regular expressions* (if you know them) can be used.

```

gapminder[grep("land", gapminder$country),]

##           country continent year lifeExp      pop gdpPercap
## 517       Finland    Europe 1952  66.550 4090500  6424.5191
## 518       Finland    Europe 1957  67.490 4324000  7545.4154
## 519       Finland    Europe 1962  68.750 4491443  9371.8426
## 520       Finland    Europe 1967  69.830 4605744 10921.6363
## 521       Finland    Europe 1972  70.870 4639657 14358.8759
## 522       Finland    Europe 1977  72.520 4738902 15605.4228
## 523       Finland    Europe 1982  74.550 4826933 18533.1576
## 524       Finland    Europe 1987  74.830 4931729 21141.0122
## 525       Finland    Europe 1992  75.700 5041039 20647.1650
## 526       Finland    Europe 1997  77.130 5134406 23723.9502
## 527       Finland    Europe 2002  78.370 5193039 28204.5906
## 528       Finland    Europe 2007  79.313 5238460 33207.0844
## 685       Iceland   Europe 1952  72.490 147962  7267.6884
## 686       Iceland   Europe 1957  73.470 165110  9244.0014
## 687       Iceland   Europe 1962  73.680 182053 10350.1591
## 688       Iceland   Europe 1967  73.730 198676 13319.8957
## 689       Iceland   Europe 1972  74.460 209275 15798.0636
## 690       Iceland   Europe 1977  76.110 221823 19654.9625
## 691       Iceland   Europe 1982  76.990 233997 23269.6075
## 692       Iceland   Europe 1987  77.230 244676 26923.2063
## 693       Iceland   Europe 1992  78.770 259012 25144.3920
## 694       Iceland   Europe 1997  78.950 271192 28061.0997
## 695       Iceland   Europe 2002  80.500 288030 31163.2020
## 696       Iceland   Europe 2007  81.757 301931 36180.7892
## 745       Ireland   Europe 1952  66.910 2952156  5210.2803
## 746       Ireland   Europe 1957  68.900 2878220  5599.0779
## 747       Ireland   Europe 1962  70.290 2830000  6631.5973
## 748       Ireland   Europe 1967  71.080 2900100  7655.5690
## 749       Ireland   Europe 1972  71.280 3024400  9530.7729
## 750       Ireland   Europe 1977  72.030 3271900 11150.9811
## 751       Ireland   Europe 1982  73.100 3480000 12618.3214
## 752       Ireland   Europe 1987  74.360 3539900 13872.8665
## 753       Ireland   Europe 1992  75.467 3557761 17558.8155
## 754       Ireland   Europe 1997  76.122 3667233 24521.9471
## 755       Ireland   Europe 2002  77.783 3879155 34077.0494
## 756       Ireland   Europe 2007  78.885 4109086 40675.9964
## 1081      Netherlands Europe 1952  72.130 10381988  8941.5719
## 1082      Netherlands Europe 1957  72.990 11026383 11276.1934
## 1083      Netherlands Europe 1962  73.230 11805689 12790.8496
## 1084      Netherlands Europe 1967  73.820 12596822 15363.2514
## 1085      Netherlands Europe 1972  73.750 13329874 18794.7457
## 1086      Netherlands Europe 1977  75.240 13852989 21209.0592
## 1087      Netherlands Europe 1982  76.050 14310401 21399.4605
## 1088      Netherlands Europe 1987  76.830 14665278 23651.3236
## 1089      Netherlands Europe 1992  77.420 15174244 26790.9496
## 1090      Netherlands Europe 1997  78.030 15604464 30246.1306
## 1091      Netherlands Europe 2002  78.530 16122830 33724.7578
## 1092      Netherlands Europe 2007  79.762 16570613 36797.9333
## 1093      New Zealand Oceania 1952  69.390 1994794 10556.5757
## 1094      New Zealand Oceania 1957  70.260 2229407 12247.3953

```

## 1095	New Zealand	Oceania	1962	71.240	2488550	13175.6780
## 1096	New Zealand	Oceania	1967	71.520	2728150	14463.9189
## 1097	New Zealand	Oceania	1972	71.890	2929100	16046.0373
## 1098	New Zealand	Oceania	1977	72.220	3164900	16233.7177
## 1099	New Zealand	Oceania	1982	73.840	3210650	17632.4104
## 1100	New Zealand	Oceania	1987	74.320	3317166	19007.1913
## 1101	New Zealand	Oceania	1992	76.330	3437674	18363.3249
## 1102	New Zealand	Oceania	1997	77.550	3676187	21050.4138
## 1103	New Zealand	Oceania	2002	79.110	3908037	23189.8014
## 1104	New Zealand	Oceania	2007	80.204	4115771	25185.0091
## 1225	Poland	Europe	1952	61.310	25730551	4029.3297
## 1226	Poland	Europe	1957	65.770	28235346	4734.2530
## 1227	Poland	Europe	1962	67.640	30329617	5338.7521
## 1228	Poland	Europe	1967	69.610	31785378	6557.1528
## 1229	Poland	Europe	1972	70.850	33039545	8006.5070
## 1230	Poland	Europe	1977	70.670	34621254	9508.1415
## 1231	Poland	Europe	1982	71.320	36227381	8451.5310
## 1232	Poland	Europe	1987	70.980	37740710	9082.3512
## 1233	Poland	Europe	1992	70.990	38370697	7738.8812
## 1234	Poland	Europe	1997	72.750	38654957	10159.5837
## 1235	Poland	Europe	2002	74.670	38625976	12002.2391
## 1236	Poland	Europe	2007	75.563	38518241	15389.9247
## 1453	Swaziland	Africa	1952	41.407	290243	1148.3766
## 1454	Swaziland	Africa	1957	43.424	326741	1244.7084
## 1455	Swaziland	Africa	1962	44.992	370006	1856.1821
## 1456	Swaziland	Africa	1967	46.633	420690	2613.1017
## 1457	Swaziland	Africa	1972	49.552	480105	3364.8366
## 1458	Swaziland	Africa	1977	52.537	551425	3781.4106
## 1459	Swaziland	Africa	1982	55.561	649901	3895.3840
## 1460	Swaziland	Africa	1987	57.678	779348	3984.8398
## 1461	Swaziland	Africa	1992	58.474	962344	3553.0224
## 1462	Swaziland	Africa	1997	54.289	1054486	3876.7685
## 1463	Swaziland	Africa	2002	43.869	1130269	4128.1169
## 1464	Swaziland	Africa	2007	39.613	1133066	4513.4806
## 1477	Switzerland	Europe	1952	69.620	4815000	14734.2327
## 1478	Switzerland	Europe	1957	70.560	5126000	17909.4897
## 1479	Switzerland	Europe	1962	71.320	5666000	20431.0927
## 1480	Switzerland	Europe	1967	72.770	6063000	22966.1443
## 1481	Switzerland	Europe	1972	73.780	6401400	27195.1130
## 1482	Switzerland	Europe	1977	75.390	6316424	26982.2905
## 1483	Switzerland	Europe	1982	76.210	6468126	28397.7151
## 1484	Switzerland	Europe	1987	77.410	6649942	30281.7046
## 1485	Switzerland	Europe	1992	78.030	6995447	31871.5303
## 1486	Switzerland	Europe	1997	79.370	7193761	32135.3230
## 1487	Switzerland	Europe	2002	80.620	7361757	34480.9577
## 1488	Switzerland	Europe	2007	81.701	7554661	37506.4191
## 1525	Thailand	Asia	1952	50.848	21289402	757.7974
## 1526	Thailand	Asia	1957	53.630	25041917	793.5774
## 1527	Thailand	Asia	1962	56.061	29263397	1002.1992
## 1528	Thailand	Asia	1967	58.285	34024249	1295.4607
## 1529	Thailand	Asia	1972	60.405	39276153	1524.3589
## 1530	Thailand	Asia	1977	62.494	44148285	1961.2246
## 1531	Thailand	Asia	1982	64.597	48827160	2393.2198
## 1532	Thailand	Asia	1987	66.084	52910342	2982.6538

```

## 1533 Thailand Asia 1992 67.298 566667095 4616.8965
## 1534 Thailand Asia 1997 67.521 60216677 5852.6255
## 1535 Thailand Asia 2002 68.564 62806748 5913.1875
## 1536 Thailand Asia 2007 70.616 65068149 7458.3963

```

There are a couple of ways of testing for more than one text value. The first uses an *or* | statement. i.e. testing if the value of `country` is `Zambia` *or* the value is `Zimbabwe`.

The `%in%` function is a convenient function for testing which items in a vector correspond to a defined set of values.

```
gapminder[gapminder$country == "Zambia" | gapminder$country == "Zimbabwe",]
```

```

##      country continent year lifeExp      pop gdpPercap
## 1681 Zambia    Africa 1952  42.038 2672000 1147.3888
## 1682 Zambia    Africa 1957  44.077 3016000 1311.9568
## 1683 Zambia    Africa 1962  46.023 3421000 1452.7258
## 1684 Zambia    Africa 1967  47.768 3900000 1777.0773
## 1685 Zambia    Africa 1972  50.107 4506497 1773.4983
## 1686 Zambia    Africa 1977  51.386 5216550 1588.6883
## 1687 Zambia    Africa 1982  51.821 6100407 1408.6786
## 1688 Zambia    Africa 1987  50.821 7272406 1213.3151
## 1689 Zambia    Africa 1992  46.100 8381163 1210.8846
## 1690 Zambia    Africa 1997  40.238 9417789 1071.3538
## 1691 Zambia    Africa 2002  39.193 10595811 1071.6139
## 1692 Zambia    Africa 2007  42.384 11746035 1271.2116
## 1693 Zimbabwe Africa 1952  48.451 3080907 406.8841
## 1694 Zimbabwe Africa 1957  50.469 3646340 518.7643
## 1695 Zimbabwe Africa 1962  52.358 4277736 527.2722
## 1696 Zimbabwe Africa 1967  53.995 4995432 569.7951
## 1697 Zimbabwe Africa 1972  55.635 5861135 799.3622
## 1698 Zimbabwe Africa 1977  57.674 6642107 685.5877
## 1699 Zimbabwe Africa 1982  60.363 7636524 788.8550
## 1700 Zimbabwe Africa 1987  62.351 9216418 706.1573
## 1701 Zimbabwe Africa 1992  60.377 10704340 693.4208
## 1702 Zimbabwe Africa 1997  46.809 11404948 792.4500
## 1703 Zimbabwe Africa 2002  39.989 11926563 672.0386
## 1704 Zimbabwe Africa 2007  43.487 12311143 469.7093

```

```
gapminder[gapminder$country %in% c("Zambia", "Zimbabwe"),]
```

```

##      country continent year lifeExp      pop gdpPercap
## 1681 Zambia    Africa 1952  42.038 2672000 1147.3888
## 1682 Zambia    Africa 1957  44.077 3016000 1311.9568
## 1683 Zambia    Africa 1962  46.023 3421000 1452.7258
## 1684 Zambia    Africa 1967  47.768 3900000 1777.0773
## 1685 Zambia    Africa 1972  50.107 4506497 1773.4983
## 1686 Zambia    Africa 1977  51.386 5216550 1588.6883
## 1687 Zambia    Africa 1982  51.821 6100407 1408.6786
## 1688 Zambia    Africa 1987  50.821 7272406 1213.3151
## 1689 Zambia    Africa 1992  46.100 8381163 1210.8846
## 1690 Zambia    Africa 1997  40.238 9417789 1071.3538
## 1691 Zambia    Africa 2002  39.193 10595811 1071.6139

```

```

## 1692 Zambia    Africa 2007 42.384 11746035 1271.2116
## 1693 Zimbabwe Africa 1952 48.451 3080907 406.8841
## 1694 Zimbabwe Africa 1957 50.469 3646340 518.7643
## 1695 Zimbabwe Africa 1962 52.358 4277736 527.2722
## 1696 Zimbabwe Africa 1967 53.995 4995432 569.7951
## 1697 Zimbabwe Africa 1972 55.635 5861135 799.3622
## 1698 Zimbabwe Africa 1977 57.674 6642107 685.5877
## 1699 Zimbabwe Africa 1982 60.363 7636524 788.8550
## 1700 Zimbabwe Africa 1987 62.351 9216418 706.1573
## 1701 Zimbabwe Africa 1992 60.377 10704340 693.4208
## 1702 Zimbabwe Africa 1997 46.809 11404948 792.4500
## 1703 Zimbabwe Africa 2002 39.989 11926563 672.0386
## 1704 Zimbabwe Africa 2007 43.487 12311143 469.7093

```

Similar to *or*, we can require that both tests are TRUE by using an *and* & operation. e.g. which years in Zambia had a life expectancy less than 40

```
gapminder[gapminder$country == "Zambia" & gapminder$lifeExp < 40,]
```

```

##       country continent year lifeExp      pop gdpPercap
## 1691 Zambia     Africa 2002 39.193 10595811 1071.614

```

Exercise

- Can you create a data frame of countries with a population less than a million in the year 2002

```

##               country continent year lifeExp      pop gdpPercap
## 95          Bahrain      Asia 2002 74.795 656397 23403.559
## 323         Comoros     Africa 2002 62.974 614382 1075.812
## 431        Djibouti     Africa 2002 53.373 447416 1908.261
## 491 Equatorial Guinea     Africa 2002 49.348 495627 7703.496
## 695         Iceland    Europe 2002 80.500 288030 31163.202
## 1019        Montenegro   Europe 2002 73.981 720230 6557.194
## 1271        Reunion     Africa 2002 75.744 743981 6316.165
## 1307 Sao Tome and Principe Africa 2002 64.337 170372 1353.092

```

- A data frame of countries with a population less than a million in the year 2002, that are not in Africa?
-
-
-

Ordering and sorting

A vector can be returned in sorted form using the `sort` function.

```
sort(countries)
sort(countries,decreasing = TRUE)
```

However, if we want to sort an entire data frame a different approach is needed. The trick is to use `order`. Rather than giving a sorted set of *values*, it will give sorted *indices*. These indices can then be used for a subset operation.

```
leastPop <- gapminder[order(gapminder$pop),]
head(leastPop)
```

```
##           country continent year lifeExp   pop gdpPercap
## 1297 Sao Tome and Principe    Africa 1952  46.471 60011 879.5836
## 1298 Sao Tome and Principe    Africa 1957  48.945 61325 860.7369
## 421      Djibouti        Africa 1952  34.812 63149 2669.5295
## 1299 Sao Tome and Principe    Africa 1962  51.893 65345 1071.5511
## 1300 Sao Tome and Principe    Africa 1967  54.425 70787 1384.8406
## 422      Djibouti        Africa 1957  37.328 71851 2864.9691
```

A final point on data frames is that we can export them out of R once we have done our data processing.

```
byWealth <- gapminder[order(gapminder$gdpPercap,decreasing = TRUE),]
head(byWealth)
```

```
##           country continent year lifeExp   pop gdpPercap
## 854      Kuwait        Asia 1957  58.033 212846 113523.13
## 857      Kuwait        Asia 1972  67.712 841934 109347.87
## 853      Kuwait        Asia 1952  55.565 160000 108382.35
## 855      Kuwait        Asia 1962  60.470 358266  95458.11
## 856      Kuwait        Asia 1967  64.624 575003  80894.88
## 858      Kuwait        Asia 1977  69.343 1140357 59265.48
```

```
write.csv(byWealth, file="dataOrderedByWealth.csv")
```

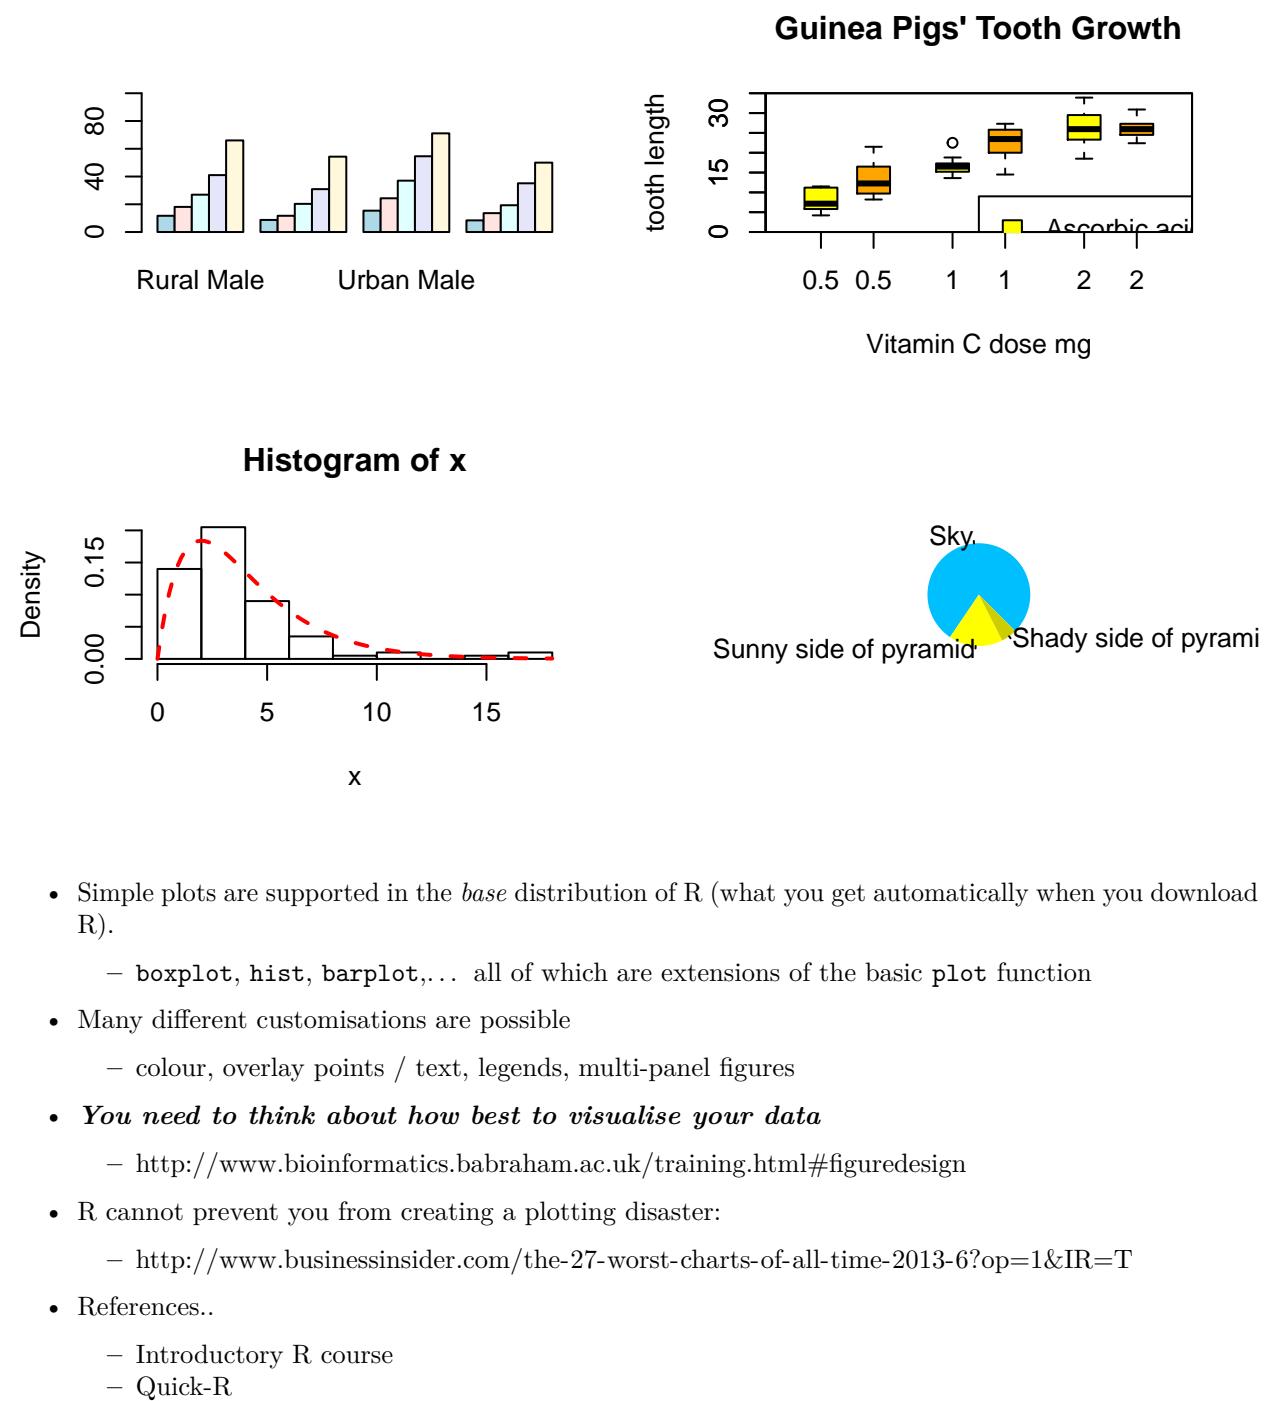
We can even order by more than one condition

```
gapminder[order(gapminder$year, gapminder$country),]
```

```
##           country continent year lifeExp   pop gdpPercap
## 1      Afghanistan     Asia 1952  28.801 8425333 779.4453
## 553      Gambia       Africa 1952  30.000 284320  485.2307
## 37      Angola        Africa 1952  30.015 4232095 3520.6103
## 1345 Sierra Leone    Africa 1952  30.331 2143249  879.7877
## 1033 Mozambique      Africa 1952  31.286 6446316  468.5260
## 193 Burkina Faso     Africa 1952  31.975 4469979  543.2552
```

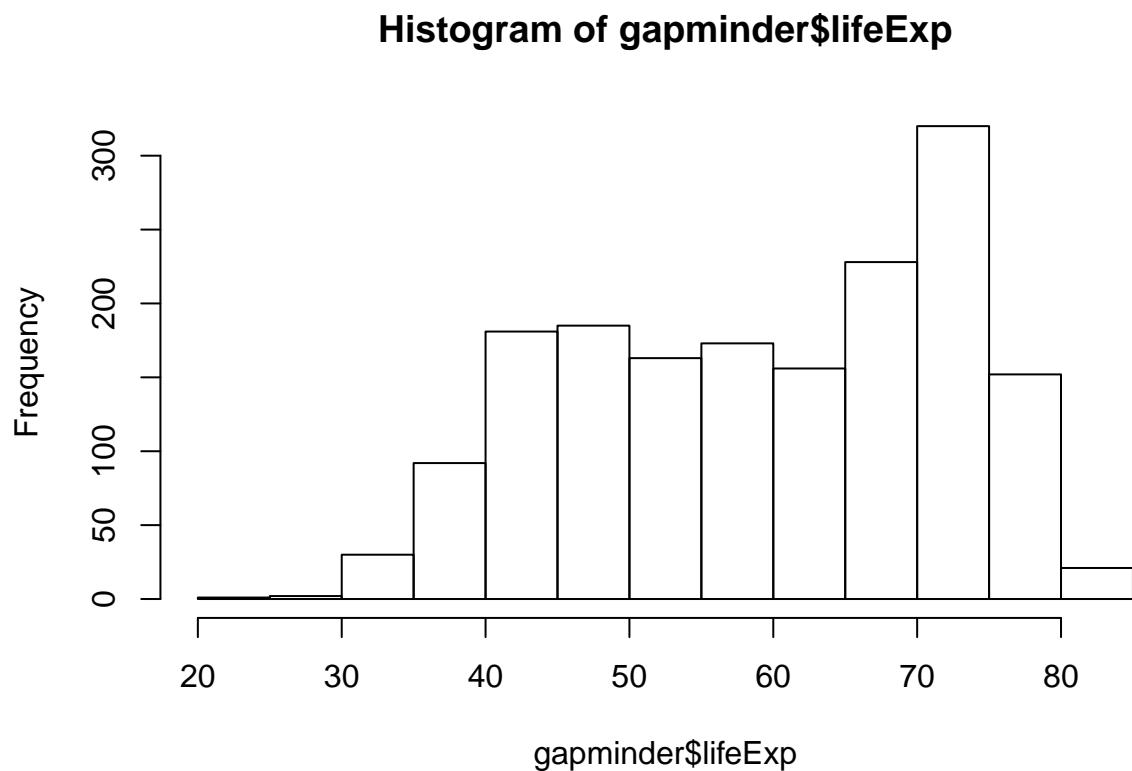
Plotting and stats (in brief!)

All your favourite types of plot can be created in R



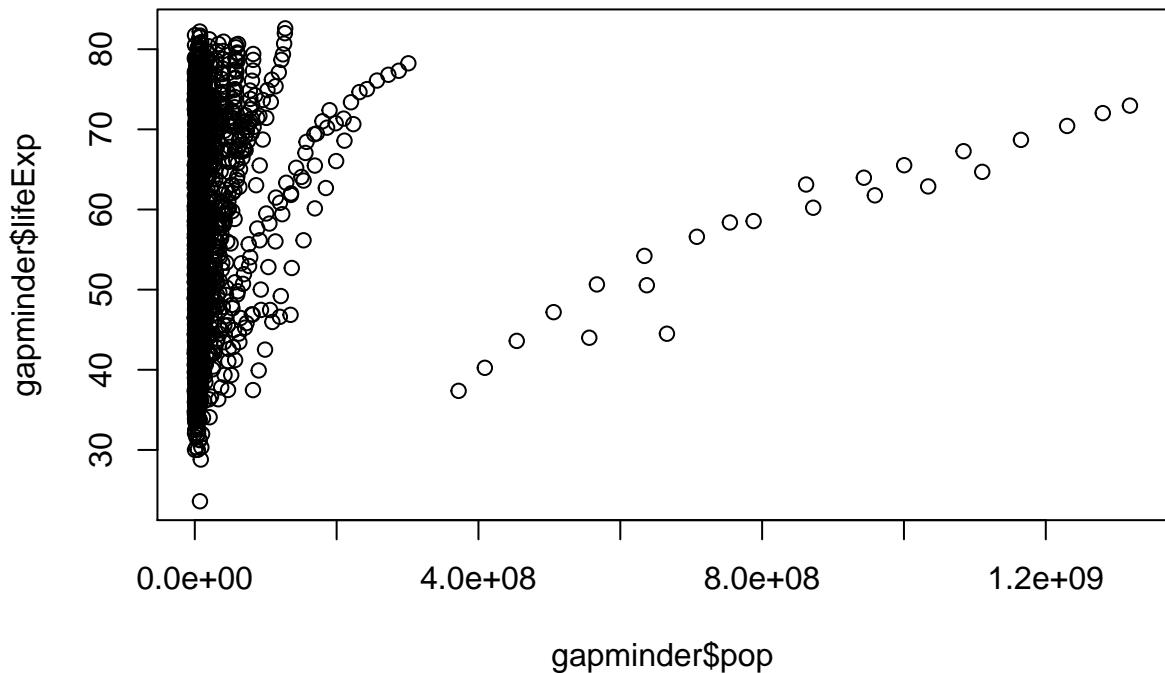
Plots can be constructed from vectors of numeric data, such as the data we get from a particular column in a data frame

```
hist(gapminder$lifeExp)
```



Scatter plots of two variables require two arguments; one for the x and one for the y axis.

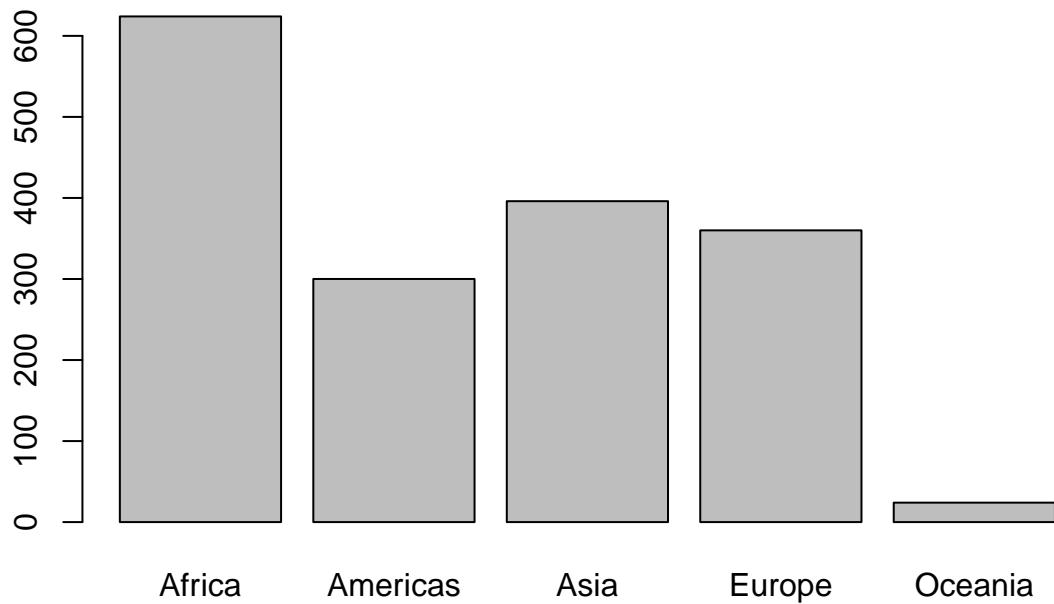
```
plot(gapminder$pop,gapminder$lifeExp)
```



What are those points on the right? We'll find out soon...

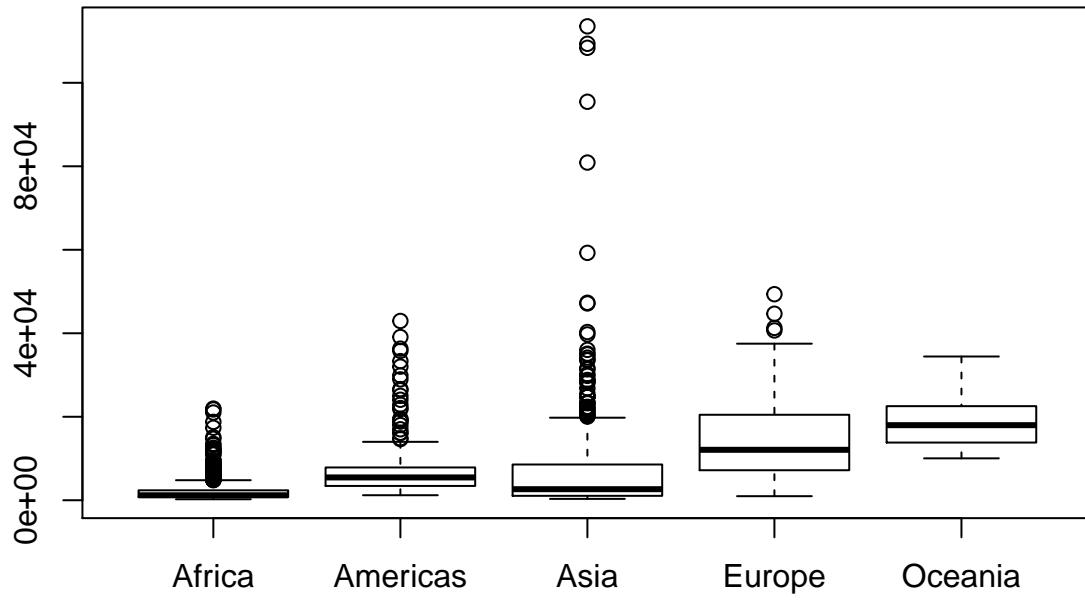
Barplots are commonly-used for counts of categorical data

```
barplot(table(gapminder$continent))
```



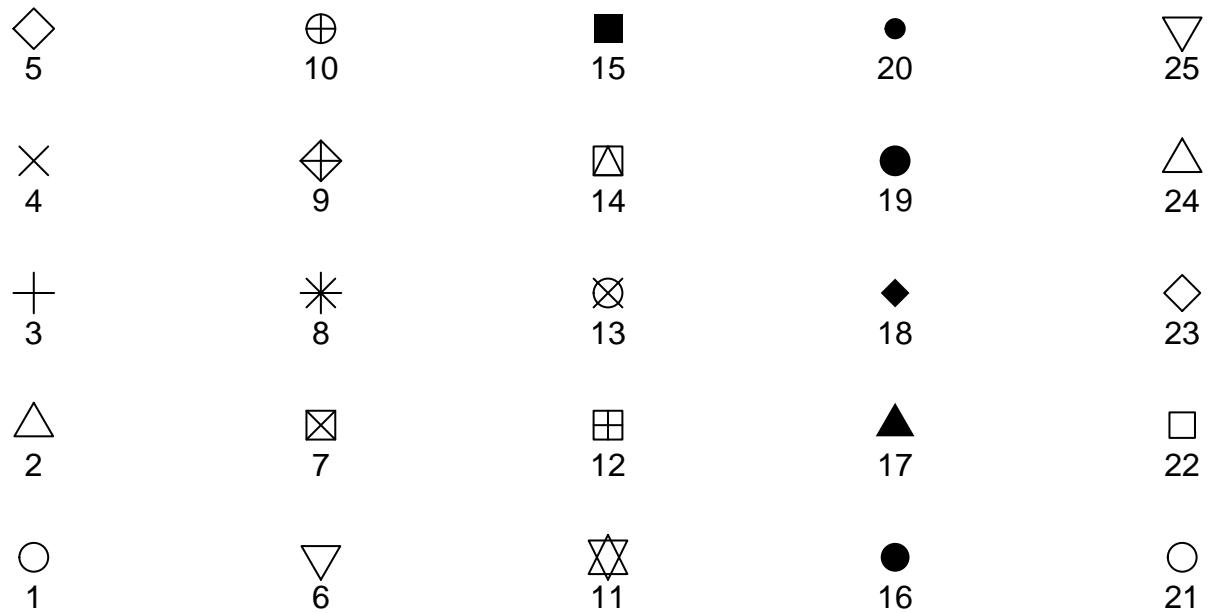
Boxplots are good for visualising and comparing distributions. Here the `~` symbol sets up a formula, the effect of which is to put the categorical variable on the x axis and continuous variable on the y axis.

```
boxplot(gapminder$gdpPerCap ~ gapminder$continent)
```



Lots of customisations are possible to enhance the appearance of our plots. Not for the faint-hearted, the help pages `?plot` and `?par` give the full details. In short,

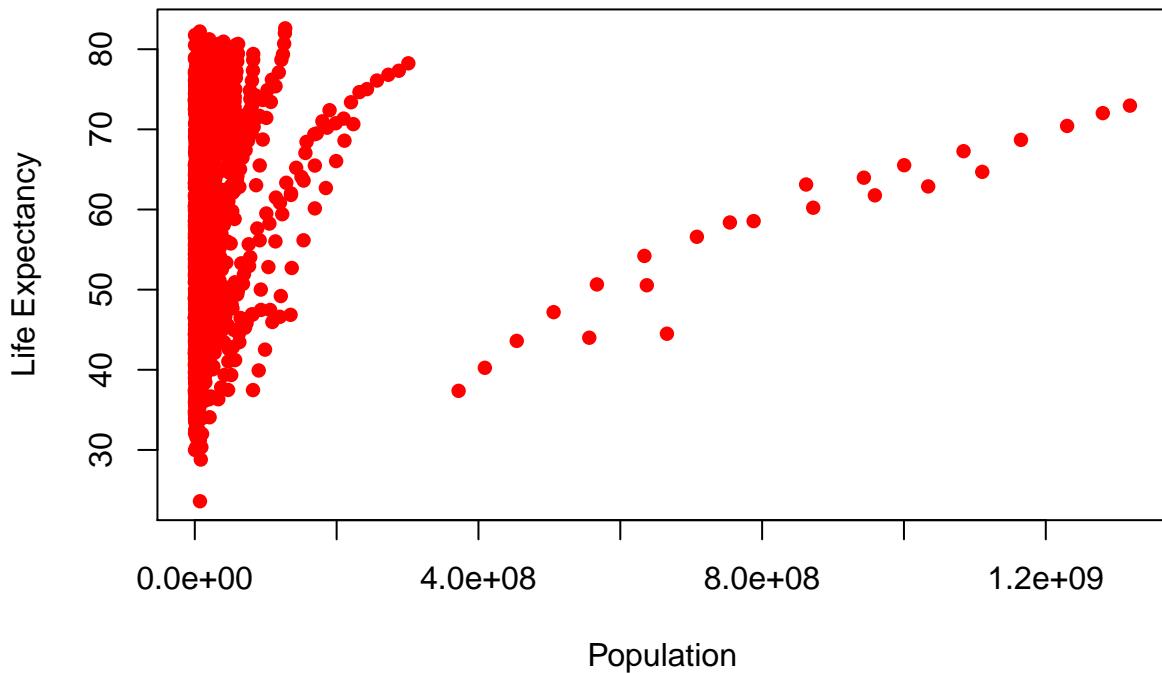
- Axis labels, and titles can be specified as character strings.
- R recognises many preset names as colours. To get a full list use `colours()`, or check this online reference.
- Plotting characters can be specified using a pre-defined number:-



Putting it all together.

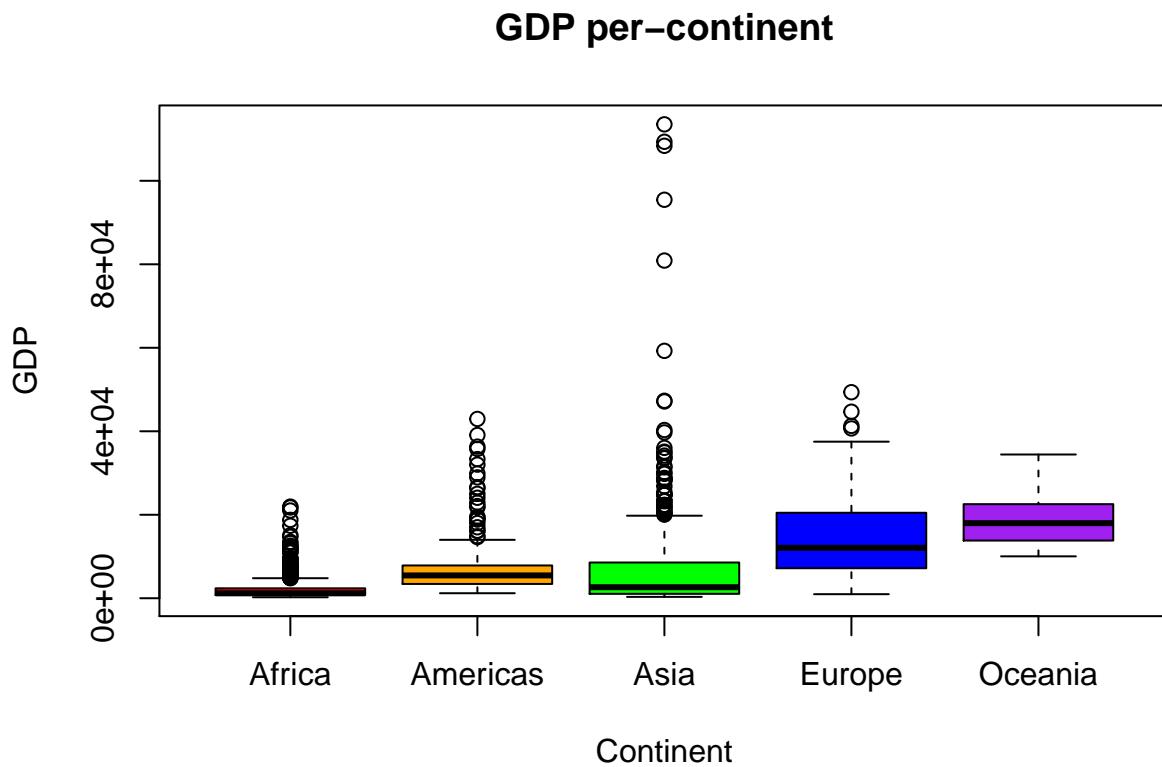
```
plot(gapminder$pop,gapminder$lifeExp,pch=16,
      col="red",ylab="Life Expectancy",
      xlab="Population",main="Life Expectancy trend with population")
```

Life Expectancy trend with population



The same customisations can be used for various plots:-

```
boxplot(gapminder$gdpPercap ~ gapminder$continent, col=c("red", "orange", "green", "blue", "purple"),
        main="GDP per-continent",
        xlab="Continent",
        ylab="GDP")
```



Plots can be exported by the ***Plots*** tab in RStudio, which is useful in an interactive setting. However, one can also save plots to a file calling the `pdf` or `png` functions before executing the code to create the plot.

```
pdf("myLittlePlot.pdf")
barplot(table(gapminder$continent))
dev.off()
```

```
## pdf
## 2
```

Any plots created in-between the `pdf(..)` and `dev.off()` lines will get saved to the named file. The `dev.off()` line is very important; without it you will not be able to view the plot you have created. `pdf` files are useful because you can create documents with multiple pages. Moreover, they can be imported into tools such as Adobe Illustrator to be incorporated with other graphics.

The canvas model

It is important to realise that base graphics in R uses a “*canvas model*” to create graphics. We can only overlay extra information on-top of an existing plot and cannot “undo” what is already drawn.

Let’s suppose we want to visualise life expectancy and population of countries in Europe and Africa. First, create two datasets to represent European and African countries in the year 2002

```
euroData <- gapminder[gapminder$continent == "Europe" & gapminder$year == 2002,]
dim(euroData)
```

```

## [1] 30 6

afrData <- gapminder[gapminder$continent == "Africa" & gapminder$year == 2002,]
dim(afrData)

```

```

## [1] 52 6

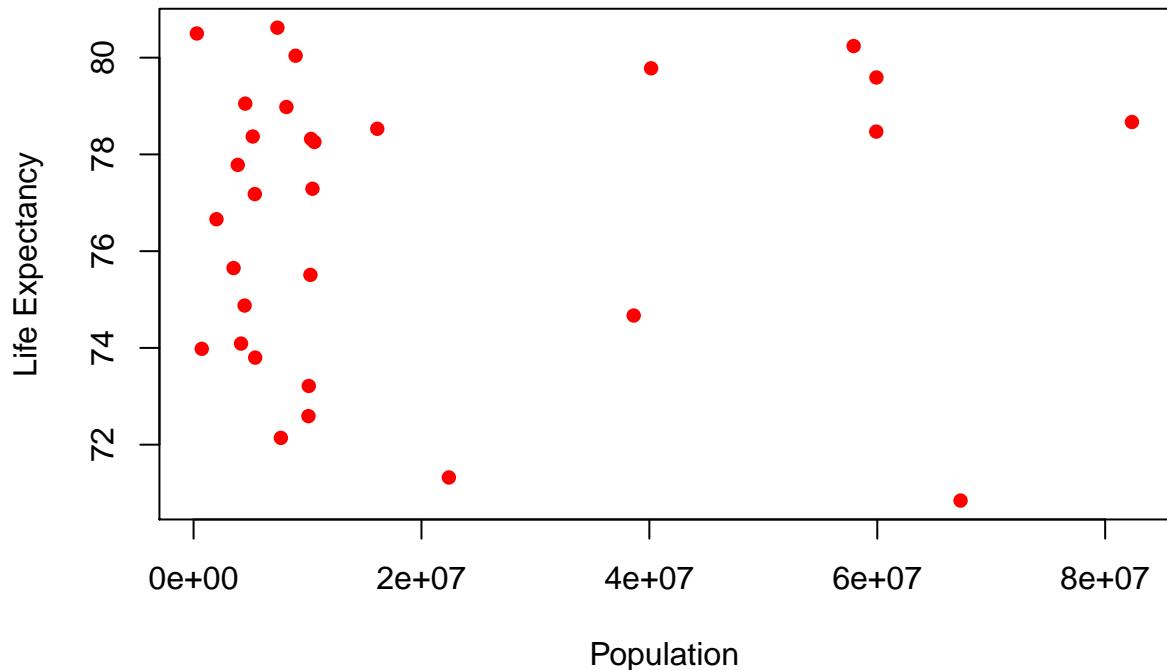
```

We can start by plotting the life expectancy of the European countries as red dots.

```

plot(euroData$pop, euroData$lifeExp, col="red",
      pch=16,
      xlab="Population",
      ylab="Life Expectancy")

```

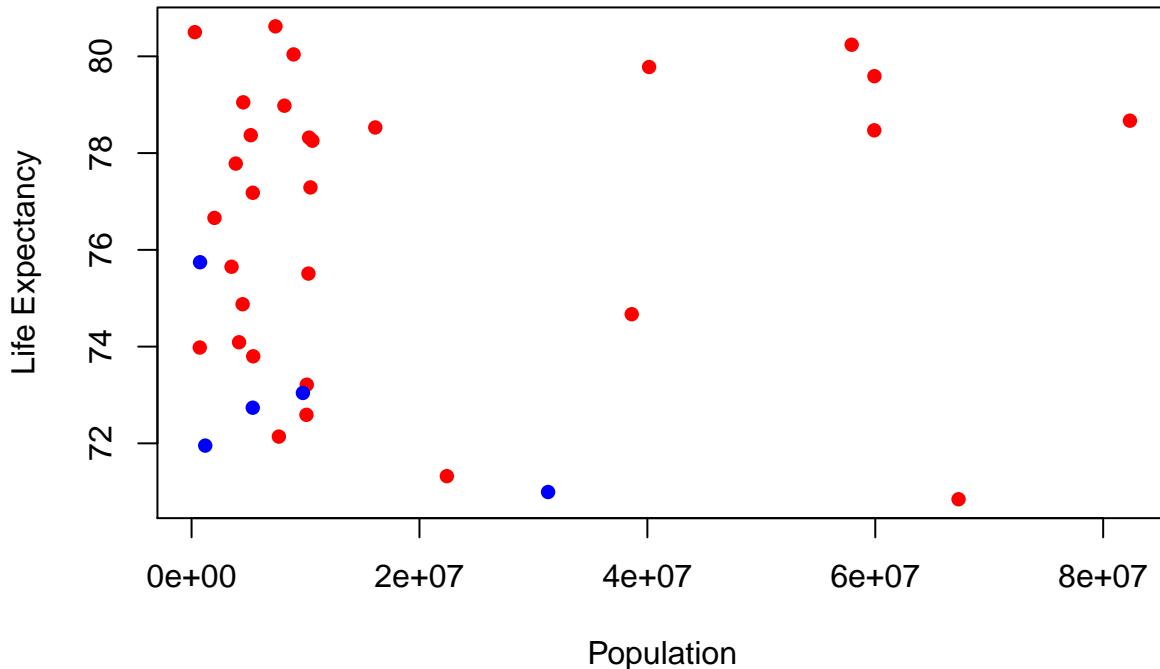


The `points` function can be used put extra points corresponding to African countries on the existing plot.

```

points(afrData$pop, afrData$lifeExp, col="blue", pch=16)

```



Wait, how many African countries did we have?

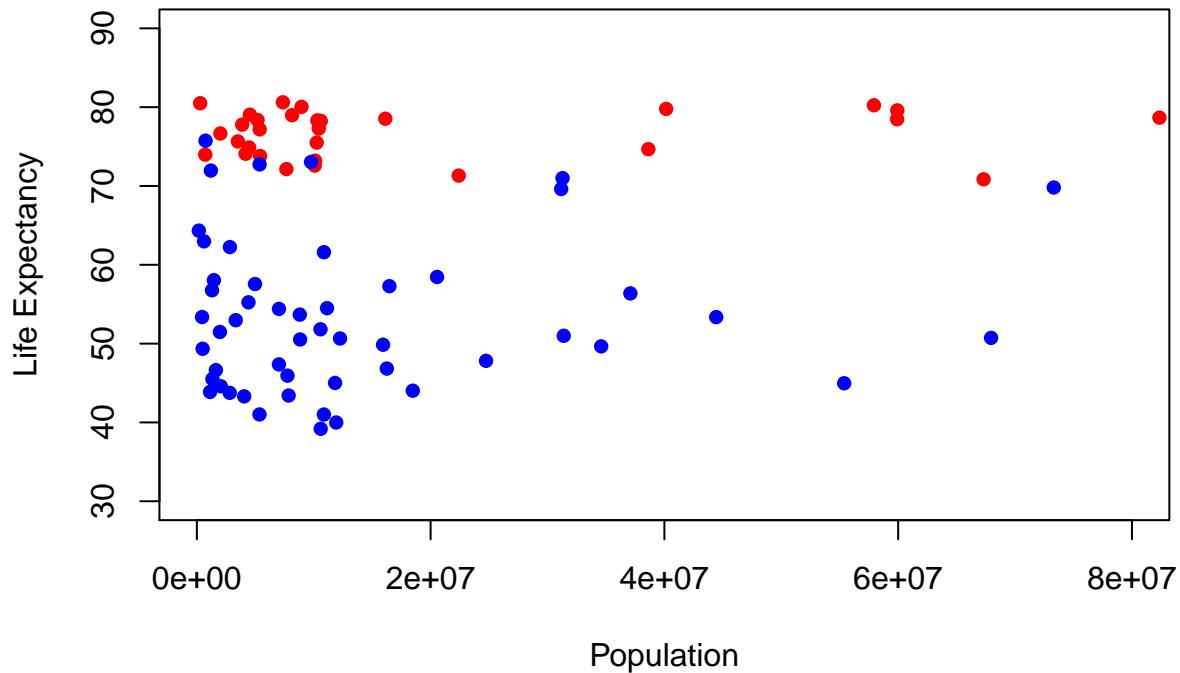
```
nrow(afrData)
```

```
## [1] 52
```

The problem here is that the initial limits of the y axis were defined using the life expectancy range of the European data. We can only add points to the existing plotting window, so any African countries with life expectancy outside this range will not get displayed.

We can define the axes when we create the plot using `xlim` and `ylim`.

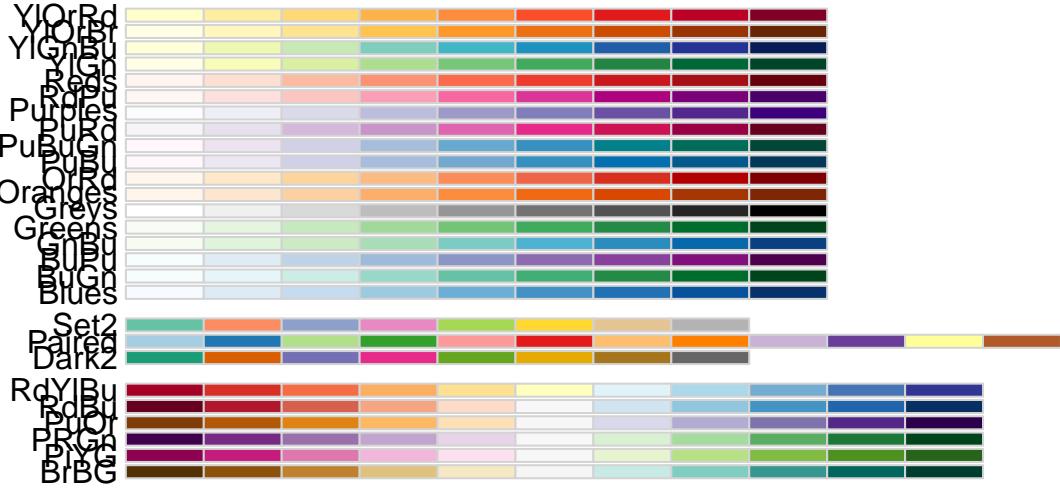
```
plot(euroData$pop, euroData$lifeExp, col="red",
      pch=16,
      xlab="Population",
      ylab="Life Expectancy",
      xlim=c(0,8e7), ylim=c(30,90))
points(afrData$pop, afrData$lifeExp, col="blue", pch=16)
```



Other useful functions for adding features to an existing plot include `text`, `abline`, `grid`, `legend` among others

Another useful trick for plotting is to take advantage of pre-existing colour palettes in R. The `RColorBrewer` package is a useful package for such palettes; many of which are friendly to those with visual impairments.

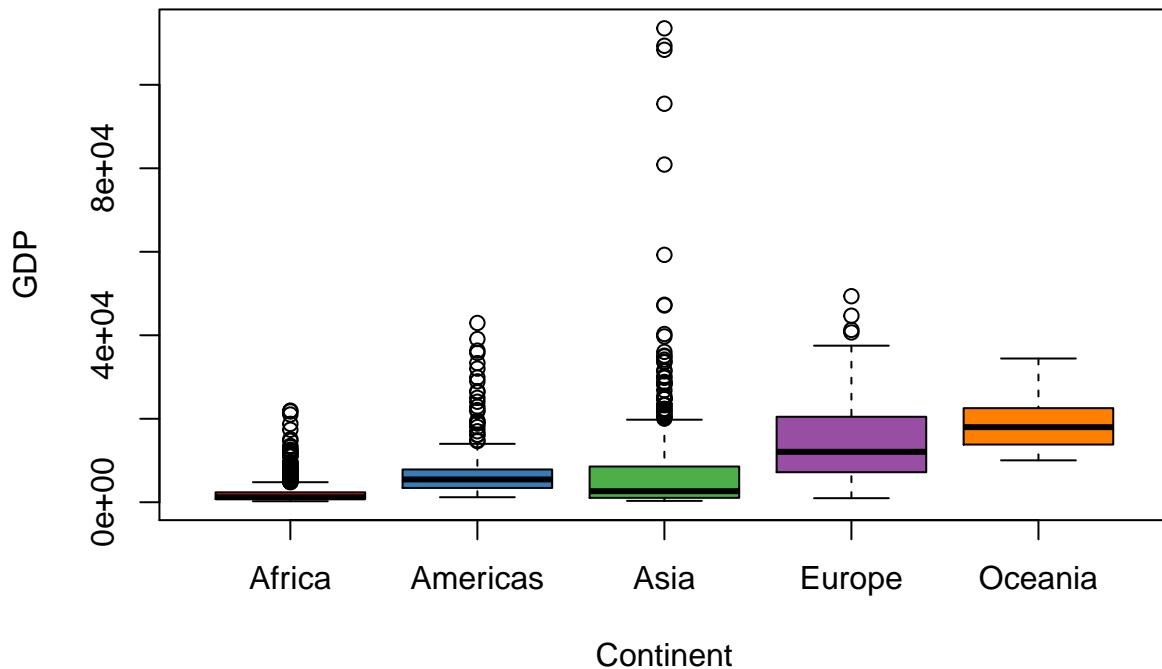
```
library(RColorBrewer)
display.brewer.all(colorblindFriendly = TRUE)
```



The `brewer.pal` function can return the names of `n` colours from one of the pre-defined palettes to be used as a `col` argument to a plotting function.

```
boxplot(gapminder$gdpPerCap ~ gapminder$continent, col=brewer.pal(5, "Set1"),
        main="GDP per-continent",
        xlab="Continent",
        ylab="GDP")
```

GDP per-continent



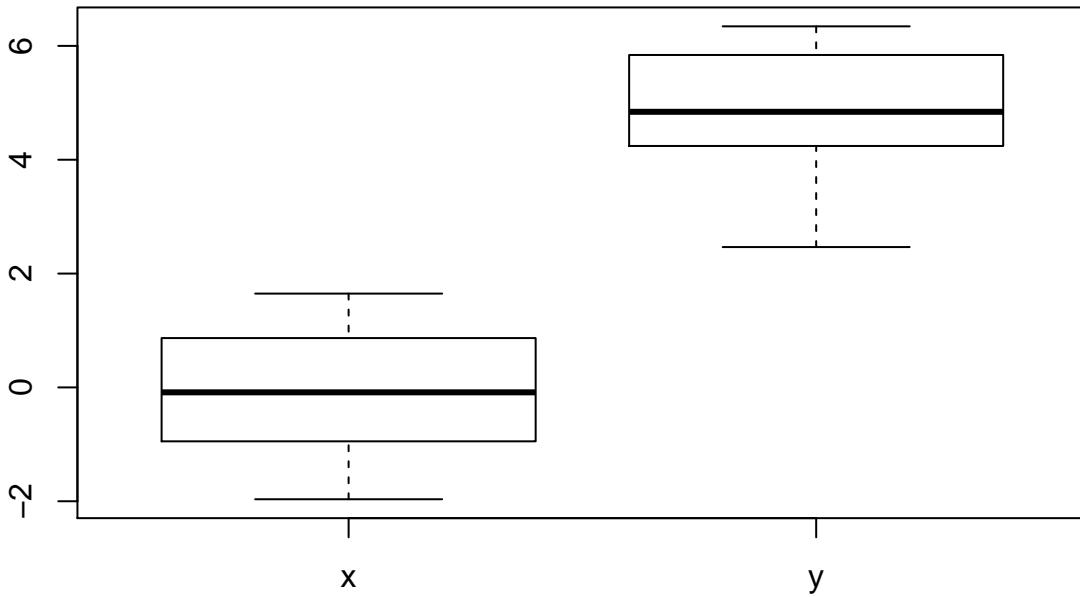
Statistical Testing

We can't really have a run-through of the R language without at least *mentioning* statistics! However, like plotting it is a vast field. The main challenges are putting your data in the correct format (which we have covered here), and deciding which test to use (**which R will not advise you on!**)

- If you have some background in statistics you can see this course from the Babraham Institute Bioinformatics Core about how to perform statistical testing in R.
- If you need a more basic grounding in which statistical test to use, you can see this course from CRUK Cambridge Institute

The `t.test` function is probably the most fundamental statistical testing function in R, and can be adapted to many different situations. Full details are given in the help page `?t.test`. Lets consider we have two vectors of normally-distributed data that we can visualise using a boxplot.

```
x <- rnorm(20)
y <- rnorm(20, 5, 1)
df <- data.frame(x,y)
boxplot(df)
```



The output from `t.test` can be used to judge if there is a statistically-significant difference in means:-

```
t.test(x,y)
```

```
##  
##  Welch Two Sample t-test  
##  
## data: x and y  
## t = -14.644, df = 37.916, p-value < 2.2e-16  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
## -5.648556 -4.276400  
## sample estimates:  
## mean of x mean of y  
## -0.1040468 4.8584312
```

If our data were paired we could set the argument `paired=TRUE` to use a different flavour of the test

```
t.test(x,y,paired = TRUE)
```

```
##  
##  Paired t-test  
##  
## data: x and y  
## t = -14.505, df = 19, p-value = 9.913e-12
```

```

## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -5.678529 -4.246427
## sample estimates:
## mean of the differences
## -4.962478

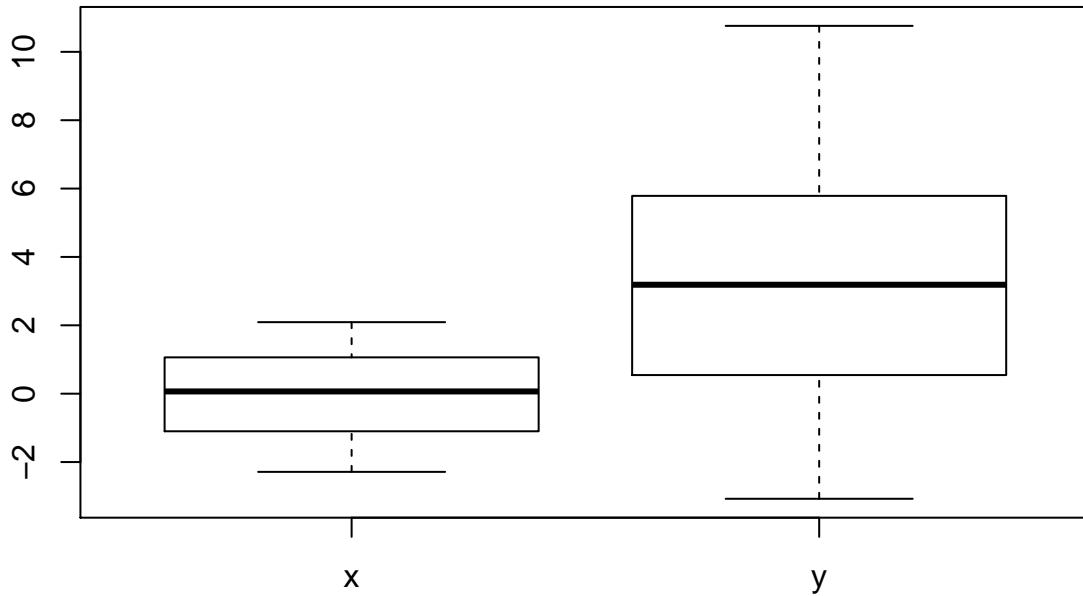
```

Similarly, if our data have different variances we can adjust the test accordingly:-

```

x <- rnorm(20)
y <- rnorm(20, 5, 4)
df <- data.frame(x,y)
boxplot(df)

```



```

t.test(x,y,var.equal = FALSE)

##
## Welch Two Sample t-test
##
## data: x and y
## t = -3.7609, df = 23.609, p-value = 0.0009824
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
## -5.067749 -1.474437
## sample estimates:

```

```
## mean of x mean of y  
## 0.04085385 3.31194675
```

Were our data not normally-distributed we could use `wilcox.test`, for example. Fortunately, most statistical tests can be accessed in a similar manner, so it is easy to switch between using different tests provided your data are in the correct format. To re-iterate, the skill is in choosing which test is appropriate.